

-----: SQL Basics Assignment Questions :-----

Q2. Explain the purpose of constraints and how they help maintain data integrity in a database. Provide examples of common types of constraints.

=> Integrity Constraints are the protocols that a table's data columns must follow. These are used to restrict the types of information that can be entered into a table. This means that the data in the database is accurate and reliable. You may apply integrity Constraints at the column or table level. The table-level Integrity constraints apply to the entire table, while the column level constraints are only applied to one column.

Here are some common types of constraints:

- **Primary key**

Uniquely identifies each record in a table, and can be made up of one or more columns.

- **Foreign key**

Ensures that a record is not loaded into a table if the foreign key values are not found in the primary key values of the parent table.

- **UNIQUE**

Prevents duplicate values from being entered into a column. This is useful for fields like email addresses or user IDs.

- **NOT NULL**

Prevents blank or null values from being entered into a column.

- **Check**

Specifies a range of rules for the contents of a table. A search condition must be satisfied for all rows in the table

Q3. Why would you apply the NOT NULL constraint to a column? Can a primary key contain NULL values? Justify your answer.

=> Why would you apply the NOT NULL constraint to a column?

= In SQL, the NOT NULL constraint in a column means that the column cannot store NULL values.

Can a primary key contain NULL values?

= The primary key cannot have NULL and duplicate values. It is used to add integrity to the table. In the case of a primary key, both Duplicate and NULL values are not valid.

Q4. Explain the steps and SQL commands used to add or remove constraints on an existing table. Provide an example for both adding and removing a constraint.

=> The **ALTER TABLE statement in SQL** is used to add, remove, or modify columns in an existing table. The ALTER TABLE statement is also used to add and remove various constraints on existing tables.

ADD – To add a new column to the table:

```
ALTER TABLE table_name  
ADD column_name datatype;
```

DROP – To delete an existing column from the table:

ALTER TABLE table_name
DROP COLUMN column_name;

Q5. Explain the consequences of attempting to insert, update, or delete data in a way that violates constraints. Provide an example of an error message that might occur when violating a constraint.

=> There are mainly three operations that have the ability to change the state of relations, these modifications are given below:

1. **Insert –**
To insert new tuples in a relation in the database.
2. **Delete –**
To delete some of the existing relation on the database.
3. **Update (Modify) –**
To make changes in the value of some existing tuples.

The ORA-00001 error message is indicative of this pattern of thinking. The ORA-00001 message is triggered when a unique constraint has been violated. Essentially the user causes the error when trying to execute an INSERT or UPDATE statement that has generated a duplicate value in a restricted field.

SQL COMMANDS :-----

Q1. Identify the primary keys and foreign keys in maven movies db. Discuss the differences.

=> **Primary vs. Foreign Keys – Key Differences**

Aspect	Primary Key	Foreign Key
Definition	Uniquely identifies a record within its own table.	Establishes a relationship between two tables.
Uniqueness	Must have unique values for each row.	Can have duplicate values across rows.
Null Values	Cannot contain NULL values.	Can contain NULL values, depending on the relationship.
Purpose	Ensures that each record in the table is unique.	Ensures data integrity by linking records across tables.
Example (in Movies DB)	MovieID in Movies, ActorID in Actors	MovieID in Movie_Actors, ActorID in Movie_Actors, DirectorID in Movie_Directors
Enforcement	Ensures each record can be uniquely identified.	Ensures that a value in one table matches a value in another (or is NULL).

Normalisation & CTE :-----

Q1. . First Normal Form (1NF):

a. Identify a table in the Sakila database that violates 1NF. Explain how you would normalize it to achieve 1NF.

=> To bring this structure into **1NF**, we need to eliminate repeating groups and ensure each field contains only a single value.

- **Option 1:** Instead of storing multiple actors' names or IDs in a single column, create a new table to represent the relationship between films and actors. For example, we could create a **film_actor** junction table with the following columns:
 - film_id (foreign key referencing film.film_id)
 - actor_id (foreign key referencing actor.actor_id)

The new structure would look like this:

- film table: Only includes unique details for each film (e.g., film_id, title, etc.).
- actor table: Contains details about each actor (e.g., actor_id, name).
- film_actor table: Contains combinations of film_id and actor_id, representing the many-to-many relationship between films and actors.

Q2. Second Normal Form (2NF):

a. Choose a table in Sakila and describe how you would determine whether it is in 2NF. If it violates 2NF, explain the steps to normalize it.

=> **Normalization to 2NF:**

To bring the table into **2NF**, we must remove the partial dependency. The steps to do so are as follows:

1. **Create a new table for customer information** to remove the dependency between last_name and first_name from the rental table.
 - We will create a **customer** table containing customer_id, last_name, and first_name.
2. **Modify the rental table:**
 - The rental table should now only store rental-specific information, which will include rental_id, rental_date, inventory_id, customer_id (as a foreign key), return_date, and staff_id.

Explanation of Normalization:

- After the normalization, the rental table is in **2NF** because:
 - All non-key attributes (such as rental_date, inventory_id, return_date, staff_id) are now fully dependent on the primary key (rental_id).
 - The attributes last_name and first_name, which were partially dependent on customer_id, have been moved to a new **customer** table, ensuring that there are no partial dependencies remaining.

Q3. Third Normal Form (3NF):

a. Identify a table in Sakila that violates 3NF. Describe the transitive dependencies present and outline the steps to normalize the table to 3NF.

=> **Steps to Normalize the Table to 3NF:**

To bring the **staff** table into **3NF**, we need to eliminate these transitive dependencies by creating separate tables for the entities that are transitively dependent. Here's how we can proceed:

1. **Create a address table:** Since **address_id** references the **address** table, we need to ensure that all address-related data (e.g., street, city, country) is moved to the **address** table to eliminate the transitive dependency between staff and address-related attributes.

- New address table:

address_id	address	city	country_id
5	'123 Elm St.'	'New York'	1
6	'456 Oak Ave.'	'Los Angeles'	2

2. **Create a store table:** If **store_id** references a **store** table, we should move store-specific details like **store_name** and **store_address** to the **store** table, and only store the foreign key (**store_id**) in the **staff** table.

- New store table:

store_id	store_name	store_address
1	'Main Store'	'123 Broadway'
2	'Sub Store'	'456 Market St'

3. **Modify the staff table:** After creating the address and store tables, the **staff** table should only contain references to those tables via foreign keys and staff-specific information (e.g., **first_name**, **last_name**, **email**, etc.).

- New staff table:

staff_id	first_name	last_name	address_id	email	store_id	active	username	password
1	'Mike'	'Hughes'	5	'mike@example.com'	1	1	'mikeh'	'pass1'
2	'Sarah'	'Williams'	6	'sarah@example.com'	1	1	'sarahw'	'pass2'

Now, **staff** no longer has any transitive dependencies. All address-related information is stored in the address table, and store-specific information is stored in the store table.